

Diagnosing a Low-Displacement Failure Mode in MAPPO for Multi-Robot Area Coverage

Harsh Pandhe

September 13, 2026

Abstract

Problem. Multi-agent reinforcement learning (MARL) is often proposed as a route to learned coordination for multi-robot exploration tasks where a classical decentralized planner is already a strong baseline, but negative results – where a MARL policy is trained under a realistic budget and simply underperforms – are rarely reported with the same rigor as positive results. **Methodology.** We train a centralized-critic MAPPO policy (Ray RLlib) on a 3-robot area-coverage task in a single simulated environment (a furnished café, ROS 2/Gazebo) for 45 training iterations, and compare it against a non-learned Frontier Heuristic controller (Voronoi partitioning, A* routing, behavior-tree mission control) and a Random Walk baseline, all filtered through an identical Control Barrier Function (CBF) safety layer. We additionally evaluate zero-shot transfer of the café-trained policy to four structurally different, unseen environments. **Key finding.** Under a 10-episode-per-condition evaluation on the training environment, the Frontier Heuristic (38.6% median area coverage rate, ACR) and even Random Walk (29.6%) outperform the trained MAPPO policy (14.5%), with the policy averaging 1.5 m of displacement per robot over a 150-step episode versus 5.4–7.4 m for the baselines. Under a separate, single-seed evaluation across all five environments (150-step episodes, $N=3$), the Frontier Heuristic outperforms MAPPO in every environment tested, by margins of $1.1\times$ to $2.0\times$. **Diagnosis.** The policy’s behavior is consistent with convergence to a low-displacement, low-risk equilibrium under a reward combining dense per-step collision penalties with a comparatively sparse cell-discovery bonus; we report the exact reward coefficients and argue this is a plausible mechanism, while explicitly noting that we did not run a reward ablation and cannot yet rule out undertraining as a contributing factor. **Contribution.** We report this negative result in full, including two independent infrastructure faults encountered and fixed during training, and argue that a narrower, evidence-scoped claim – that MAPPO exhibited a reproducible low-displacement failure mode and did not match a strong classical planner under the tested reward and training budget – is more defensible and more useful to practitioners than a general claim that classical methods beat MARL.

1 Introduction

Multi-robot area coverage – driving a team of robots to jointly maximize explored floor area – has a strong, well-understood non-learned baseline: frontier-based exploration combined with a path planner and a decentralized space-partitioning scheme. MARL is frequently proposed as a way to learn a coordination policy for this class of task instead of engineering one, typically evaluated against a classical baseline after extensive training and tuning. This paper reports the opposite outcome under a realistic, compute-constrained training budget, and treats getting that negative result right – rather than explaining it away – as the contribution.

We deliberately narrow the claim to what our evidence supports: this is a report of a specific, reproducible failure mode of one MAPPO configuration on one task family, not a general claim

about MARL versus classical planning. Section 2 details the system, reward, and two distinct evaluation protocols we ran (they use different sample sizes and are not numerically interchangeable). Section 3 reports both. Section 4 discusses the mechanism and its evidentiary limits.

2 System and Method

2.1 Platform and Safety Layer

Three TurtleBot3 Waffle robots (differential drive, 2D LiDAR, IMU) are simulated in Gazebo Harmonic under ROS 2 Jazzy [3, 4], coordinated via a PettingZoo-compatible parallel multi-agent environment [8]. Every commanded velocity, regardless of which controller produced it, passes through an identical unicycle Control Barrier Function (CBF) filter [1], solved as a quadratic program via OSQP [7] at each control step, with separate safety radii for static obstacles ($d_{\text{safe,obs}} = 0.22\text{ m}$) and neighboring agents ($d_{\text{safe,agent}} = 0.45\text{ m}$). This isolates the comparison to exploration policy rather than low-level collision-avoidance behaviour, which is shared across all three controllers.

2.2 Controllers Compared

Frontier Heuristic: a non-learned controller combining frontier cell selection, 8-connected grid A* [2] waypoint routing with obstacle inflation, unreachable-target blacklisting to break wall-lock deadlocks, and a `py_trees` behavior tree for mission-level state (explore, recover-from-stuck, patrol-on-saturation). This controller encodes substantial task-specific structural knowledge (a discretized map, a graph search algorithm, an explicit stuck-recovery policy) that a learned policy must instead acquire implicitly from reward and observations; we treat this asymmetry as part of the research question (when is engineering that structure worthwhile versus learning it?) rather than an oversight.

Random Walk: uniform random heading changes at fixed intervals, CBF-filtered identically to the other two controllers.

MAPPO: a centralized-critic, shared-policy PPO [6] implementation via Ray RLlib [5]. Each robot observes a 46-dimensional vector: 24 LiDAR range bins, its own pose and velocity, and its assigned goal’s relative position; the policy is a 2-layer MLP (256 units, ReLU) shared across agents, with a separate centralized critic that additionally observes all agents’ states. The action space is continuous, $a_i = [v_i, \omega_i]$ (linear and angular velocity), passed through the shared CBF filter described above. Training used PPO clip 0.2, $\gamma = 0.99$, GAE $\lambda = 0.95$, learning rate 10^{-4} , entropy coefficient 0.01, train batch size 700, minibatch size 64, 5 SGD epochs per iteration, for 45 iterations (approximately 31,500 environment transitions in total). We report this total transition count explicitly because “45 iterations” alone does not convey the sample budget: this is on the order of 3×10^4 environment steps, at the low end of what is typically reported for PPO-family algorithms on continuous-control tasks, and we cannot rule out that additional training would change the outcome (see Section 4).

2.3 Reward Function

The per-robot, per-step reward is

$$r_t = -0.01 + 5.0 \Delta d_{\text{goal}} + 4.0 n_{\text{new-cells}} - 8.0 \mathbb{I}[\text{wall collision}] - p_{\text{proximity}} + 20 \mathbb{I}[\text{goal reached}] \quad (1)$$

where Δd_{goal} is the reduction in distance to the robot’s current goal since the previous step, $n_{\text{new-cells}}$ is the number of previously-unvisited grid cells the swarm collectively discovered this step, and

$p_{\text{proximity}}$ is 0.5 if any other agent is within 0.5 m and an additional 10.0 if within 0.20 m (counted as a collision after the first 15 steps of an episode). The per-cell discovery bonus was raised from an earlier value of 2.0 to 4.0 specifically because it was judged too small relative to the ± 20 goal/collision terms to provide useful gradient signal toward the coverage behaviour the policy is evaluated on; even after this change, a single wall collision (-8.0) or close-proximity event (-10.0) is substantially larger in magnitude than the reward for discovering several new cells in a single step (4.0 each), which is the quantitative basis for our hypothesis in Section 4.

2.4 Area Coverage Rate (ACR)

We define

$$\text{ACR} = \frac{N_{\text{visited reachable cells}}}{N_{\text{reachable cells}}} \times 100 \quad (2)$$

computed on a fixed-resolution (0.4 m) occupancy grid per environment. A cell is *reachable* if it is not classified as a static obstacle by the environment’s known obstacle map; cells occupied only transiently by a robot are not excluded from the denominator. A cell counts as *visited* the first time any robot’s footprint overlaps it; repeat visits do not increase ACR but are separately tracked as *overlap redundancy* (mean visits per visited cell), reported in Figure 1.

2.5 Environments and Train/Test Split

Five Gazebo worlds are used: **cafe** (furnished, 17×8 m, 800 grid cells), **warehouse** (open logistics hall, 30×50 m, 9,375 cells), **depot** (industrial, structural pillars and boxsets, 29×15 m, 3,000 cells), **office** (partitioned floor plan, narrow doorways, 26.2×18.3 m, 3,500 cells), and **maze** (labyrinthine corridors, 64×11 m, 4,800 cells, of which only $\approx 17\%$ is physically reachable floor space). **The MAPPO policy is trained exclusively on cafe.** Evaluation on **warehouse**, **depot**, **office**, and **maze** is therefore zero-shot transfer, not in-distribution evaluation, and should be read as a generalization test rather than a like-for-like comparison of learning capability per environment; we did not train separate policies per environment, which we identify as the most valuable follow-up experiment in Section 4.

2.6 Two Evaluation Protocols

We report results from two separate evaluation runs that differ in sample size and (potentially) exact policy checkpoint, and we do not treat their numbers as directly interchangeable:

Protocol A (café deep-dive): 10 episodes per condition (Random Walk, Frontier Heuristic, MAPPO nominal, MAPPO with injected Gaussian LiDAR noise, MAPPO with single-agent failure), 150-step episodes, in **cafe** only. This is the source of the 14.5%/38.6%/ 29.6% ACR figures and the 1.5 m displacement figure in the abstract, and of Figure 1.

Protocol B (five-world matrix): a single seed (42), one episode per (world, controller) pair, 150-step episodes, across all five environments and three controllers (plus two MAPPO perturbation scenarios), reported in Table 1. This protocol trades per-condition replication for environment breadth.

Because Protocol B uses $n=1$ per condition, we do not report confidence intervals for it and treat it as indicative rather than statistically conclusive; its value is in showing the same qualitative ranking (heuristic > MAPPO) holds across five structurally different environments, not in the precision of any single number. Protocol A, with $n=10$ per condition, is our primary quantitative evidence and is reported with the full distribution in Figure 1.

3 Results

3.1 Protocol B: Five-World Comparison

Table 1 reports ACR and total swarm displacement from the single-seed, five-world matrix.

Table 1: Protocol B: ACR (%) and total swarm displacement (m), 150-step episodes, $N=3$, single seed. Ratios in the last column are Heuristic / MAPPO.

World	Heuristic ACR	MAPPO ACR	Random ACR	Heuristic Dist.	Ratio (H/M)
cafe	57.0	28.2	22.6	16.5	$2.0\times$
warehouse	4.3	3.3	2.6	31.0	$1.3\times$
depot	13.3	6.7	6.8	12.1	$2.0\times$
office	3.8	3.0	1.7	14.8	$1.3\times$
maze	2.7	2.2	1.5	10.9	$1.2\times$

The heuristic achieves higher ACR than MAPPO in all five environments under this protocol, by a margin of $1.2\times$ to $2.0\times$ – we note this range explicitly because it differs from, and should not be conflated with, the Protocol A ratio discussed next.

3.2 Protocol A: Café Distributional Comparison

Figure 1 shows the 10-episode-per-condition distribution in **café**. Under this protocol, median ACR is 38.6% (heuristic), 29.6% (Random Walk), and 14.5% (MAPPO nominal) – a heuristic/MAPPO ratio of $2.7\times$, and MAPPO trails even the unplanned Random Walk baseline. Mean per-robot displacement over the 150-step episode is 5.4 m (heuristic), 7.4 m (Random Walk), and 1.5 m (MAPPO).




Figure 1: Protocol A: area coverage rate and cell-visit overlap redundancy, 10 episodes per condition, café environment. MAPPO’s tight, low-value boxes indicate consistent underperformance across episodes, not high-variance noise.

We flag that Protocol A and Protocol B were run against what we believe are the same trained policy weights, but we cannot fully rule out checkpoint or environment-seed drift between the two

evaluation runs given how they were orchestrated; readers should treat Table 1’s $1.2\text{--}2.0\times$ range and Figure 1’s $2.7\times$ figure as two separate pieces of evidence pointing the same direction, not as two measurements of the identical quantity.




Figure 2: Frontier Heuristic coverage heatmap, café environment (visited cells in light tone, obstacles mid-tone, unexplored dark, per-agent trajectories overlaid). We were not able to generate a directly comparable MAPPO trajectory heatmap for this revision; given the policy’s near-static behaviour (Section 3), we expect it to show almost no spatial extent, and consider this a priority addition for a future version of this paper.

4 Discussion

4.1 Proposed Mechanism, and What We Have Not Shown

The policy’s near-static behaviour (1.5 m displacement over 150 steps, versus 5–10 m for both non-learned baselines) is consistent with convergence to a low-risk equilibrium in which the expected cost of attempting to move and discover new cells (a -8.0 wall-collision penalty or -10.0 close-proximity penalty, either of which can occur from a single misjudged action) exceeds the expected benefit (4.0 per newly discovered cell, requiring successful, sustained movement to accumulate). We consider this a plausible, quantitatively grounded hypothesis given the reward coefficients in Section 2, but we did not run a reward ablation (e.g., reducing the collision penalty magnitude while holding the discovery bonus fixed, or vice versa) to establish it as a demonstrated causal mechanism, and we did not run a longer training budget to rule out that the policy would eventually escape this equilibrium with more iterations. Both are the highest-value follow-up experiments we identify for future work. We therefore state our conclusion narrowly: the observed underperformance *coincides with* convergence to a low-displacement, penalty-averse policy, under a reward and training budget we have fully specified; we have not experimentally eliminated undertraining, reward-scale miscalibration, or the zero-shot generalization gap (Section 2) as contributing factors.

4.2 Scope of the Claim

We deliberately frame this as an engineering question – when should a practitioner prefer an engineered planner over a MARL policy for this class of task, under a realistic training budget – rather than a general claim that classical heuristics outperform MARL. The Frontier Heuristic compared here embeds substantial hand-engineered structure (a discretized map, A*, explicit deadlock recovery) that MAPPO must instead infer implicitly; a fair reading of our result is that this structure was worth more, under our training budget, than what the policy was able to learn in its place, not that learning cannot in principle recover it given more compute, a better reward, or per-environment training.

5 Conclusion

Under a fully specified, single-environment training budget and reward function, and evaluated both on the training environment (10 episodes per condition) and zero-shot across four additional, structurally distinct environments (single-seed matrix), a MAPPO policy exhibited a reproducible low-displacement failure mode and did not match a strong classical frontier-and-planning baseline. We report the exact reward coefficients, training budget, and both evaluation protocols in full, including their differing sample sizes and a residual uncertainty about checkpoint consistency between them, so that readers can weigh the evidence directly rather than take our summary on faith. We identify a reward-coefficient ablation, a longer training budget, and per-environment (rather than zero-shot) training as the three highest-value experiments to strengthen or overturn this result.

A Infrastructure Faults Encountered During Training

Two infrastructure faults were encountered and resolved during this work. First, a segmentation fault inside `torch.optim.Adam.__init__` under Ray actor threads was traced to a PyTorch TorchDynamo compilation hook interacting with Ray’s threading model; disabling TorchDynamo (`TORCHDYNAMO_DISABLE=1`) resolved it. Second, training halted at iteration 34/45 when a host Wi-Fi disassociation event caused Ray’s distributed control plane (bound by default to the host’s dynamic IP rather than loopback) to lose connectivity to its own coordination service; binding Ray explicitly to `127.0.0.1` and capping its object store in-process resolved this and allowed training to resume and complete to iteration 45. Neither fault altered the training configuration, reward function, or evaluation protocol reported in the main text; both are reported for reproducibility.

References

- [1] Aaron D Ames, Samuel Coogan, Magnus Egerstedt, Gennaro Notomista, Koushil Sreenath, and Paulo Tabuada. Control barrier functions: Theory and applications. *2019 18th European Control Conference (ECC)*, pages 3420–3431, 2019.
- [2] Peter E Hart, Nils J Nilsson, and Bertram Raphael. A formal basis for the heuristic determination of minimum cost paths. In *IEEE Transactions on Systems Science and Cybernetics*, volume 4, pages 100–107, 1968.
- [3] Steven Macenski, Tully Foote, Brian Gerkey, Chris Lalancette, and William Woodall. Robot operating system 2: Design, architecture, and uses in the wild. In *Science Robotics*, volume 7, 2022.

- [4] Open Robotics. Gazebo simulator. <https://gazebo.org/>, 2024.
- [5] Ray Project. RLlib: Scalable reinforcement learning. <https://docs.ray.io/en/latest/rllib/index.html>, 2024.
- [6] John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. Proximal policy optimization algorithms. *arXiv preprint arXiv:1707.06347*, 2017.
- [7] Bartolomeo Stellato, Goran Banjac, Paul Goulart, Alberto Bemporad, and Stephen Boyd. OSQP: An operator splitting solver for quadratic programs. *Mathematical Programming Computation*, 12(4):637–672, 2020.
- [8] Justin Terry, Benjamin Black, Nathaniel Grammel, Mario Jayakumar, Ananth Hari, Ryan Sullivan, Luis S Santos, Clemens Dieffendahl, Caroline Horsch, Rodrigo Perez-Vicente, et al. PettingZoo: Gym for multi-agent reinforcement learning. In *Advances in Neural Information Processing Systems*, volume 34, pages 15032–15043, 2021.